

REPORT ECE408 Milestone1

Teamname: spicy-chicken

Members: Xiaocong Yu(xy21), Dawei Wang(dwang56) , Kexuan Zou(kzou3)

List of kernel calls that collectively consume more than 90% of the program time

| Type | Time(%) | Time | Calls | Avg | Min | Max | Name |
|---|---------|----------|-------|----------|----------|----------|---|
| GPU activities: | 40.07% | 16.772ms | 20 | 838.61us | 1.1200us | 16.152ms | [CUDA memcpy HtoD] |
| | 20.15% | 8.4355ms | 1 | 8.4355ms | 8.4355ms | 8.4355ms | void |
| cudnn::detail::implicit_convolve_sgemm<float, float, int=1024, int=5, int=5, int=3, int=3, int=3, int=1, bool=1, bool=0, bool=1>(int, int, int, float const *, int, float*, cudnn::detail::implicit_convolve_sgemm<float, float, int=1024, int=5, int=5, int=3, int=3, int=3, int=1, bool=1, bool=0, bool=1>*, kernel_conv_params, int, float, float, int, float, float, int, int) | | | | | | | |
| | 11.82% | 4.9474ms | 1 | 4.9474ms | 4.9474ms | 4.9474ms | volta_cgemm_64x32_tn |
| | 7.04% | 2.9486ms | 2 | 1.4743ms | 25.087us | 2.9235ms | void |
| op_generic_tensor_kernel<int=2, float, float, float, int=256, cudnnGenericOp_t=7, cudnnNanPropagation_t=0, cudnnDimOrder_t=0, int=1>(cudnnTensorStruct, float*, cudnnTensorStruct, float const *, cudnnTensorStruct, float const *, float, float, float, float, dimArray, reducedDivisorArray) | | | | | | | |
| | 5.71% | 2.3909ms | 1 | 2.3909ms | 2.3909ms | 2.3909ms | void fft2d_c2r_32x32<float, bool=0, bool=0, unsigned int=1, bool=0, bool=0>(float*, float2 const *, int, int, int, int, int, int, int, int, float, float, cudnn::reduced_divisor, bool, float*, float*, int2, int, int) |
| | 5.59% | 2.3403ms | 1 | 2.3403ms | 2.3403ms | 2.3403ms | volta_sgemm_128x128_tn |
| | 4.56% | 1.9076ms | 1 | 1.9076ms | 1.9076ms | 1.9076ms | void |
| cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0, bool=0>(cudnnTensorStruct, float const *, cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0, bool=0>, cudnnTensorStruct*, cudnnPoolingStruct, float, cudnnPoolingStruct, int, cudnn::reduced_divisor, float) | | | | | | | |
| | 4.21% | 1.7638ms | 1 | 1.7638ms | 1.7638ms | 1.7638ms | void fft2d_r2c_32x32<float, bool=0, unsigned int=0, bool=0>(float2*, float const *, int, int, int, int, int, int, int, int, int, cudnn::reduced_divisor, bool, int2, int, int) |

List of API calls that collectively consume more than 90% of the program time

| | Type | Time(%) | Time | Calls | Avg | Min | Max | Name |
|------------|--------|----------|------|----------|----------|----------|-----|---------------------------|
| API calls: | 41.37% | 3.00690s | 22 | 136.68ms | 13.759us | 1.58874s | | cudaStreamCreateWithFlags |
| | 33.85% | 2.46087s | 24 | 102.54ms | 87.306us | 2.45588s | | cudaMemGetInfo |
| | 21.29% | 1.54779s | 19 | 81.463ms | 948ns | 413.98ms | | cudaFree |

Explanation of the difference between kernels and API calls

“Summary mode is the default operating mode for nvprof. In this mode, nvprof outputs a single result line for each kernel function and each type of CUDA memory copy/set performed by the application. For each kernel, nvprof outputs the total time of all instances of the kernel or type of memory copy as well as the average, minimum, and maximum time. The time for a kernel is the kernel execution time on the device. By default, nvprof also prints a summary of all the CUDA runtime/driver API calls. Output of nvprof (except for tables) are prefixed with ==<pid>==, <pid> being the process ID of the application being profiled.” -- CUDA Toolkit Reference

From the official reference for nvprof tool, we know that the list for kernels and API are different since the API calls are mostly executed at CPU side (host code that may or may not invoke GPU), whereas the kernel calls are executed at GPU side (device code). There list some driver API calls including cudaGetDevice, cuDeviceGetName and Runtime API calls including cudaFunctionSetAttr, cudaMemsetAsync, etc.

API calls deals with collecting information for NVCC during compile time to help generate architecture and computing-ability specific executable code for CPU and GPU, and invoking kernel function on device so it will be executed much more times than kernel calls. Whereas Kernel time usage information above is collected for single kernel, so you can see the calculation functions are only called once. Time mostly comprised of kernel calculation function and cudaMemcpy from the GPU shared memory.

Output of rai running MXNet on CPU & GPU

Running /usr/bin/time python m1.1.py

Loading fashion-mnist data... done

Loading model... done

New Inference

EvalMetric: {'accuracy': 0.8236}

9.18user 3.70system

Time used 0:05.38 elapsed 239%CPU

(0avgtext+0avgdata2470492maxresident)k0inputs+2824outputs (0major+669491minor)pagefaults

0swaps

Time used for m1.1py

0:05.38 elapsed 239%CPU

Running /usr/bin/time python m1.2.py

Loading fashion-mnist data... done

Loading model... done

New Inference

EvalMetric: {'accuracy': 0.8236}

4.31user 3.25system

0:04.23 elapsed

178%CPU

(0avgtext+0avgdata2858044maxresident)k8inputs+1728outputs (0major+663219minor)pagefaults

0swaps

Time used for m1.2py

0:04.23 elapsed 178%CPU